



CS511, FALL 2024



# Breaking Bao

## Evaluating Bao on High-Complexity Benchmarks

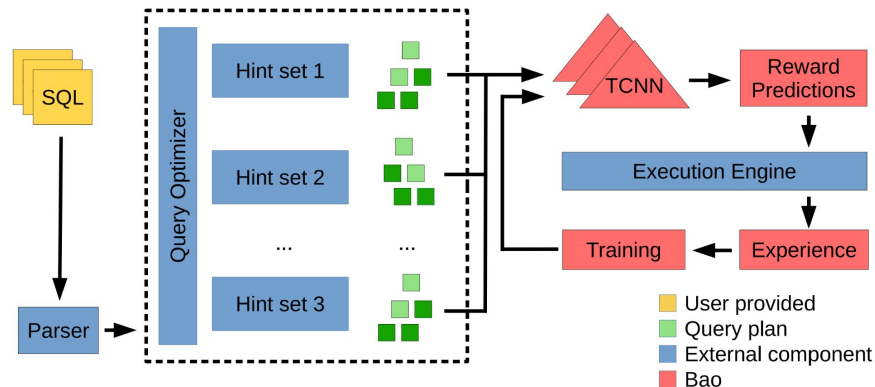
Presented by (Team20):

Vidhi Rambhia, Saharsh Barve  
Rahul Bothra, Omkar Dhekane

# Introduction to Query Optimizers

Query optimizers generate an efficient 'plan' to execute the query.

**Bao:** Uses RL to optimize PostgreSQL with query hints via reinforcement learning.



How do we know that a Query Optimizer is performant?

# Motivation



**Q: How do we know if a query optimizer is performant?**

A: Evaluate performance over a set of standard benchmarks.

**Production databases can deviate from the controlled benchmark:**

- Complex queries, evolving workloads, and schema changes.
- Examples: Shifts from selective queries to broad scans, schema updates (eg: index drops).

**Q: How do we know if a query optimizer is robust?**

A: ??

# Related Work



## Landscape of Learned Optimizers:

- **Bao:** Enhances optimizers with ML-based query hints.
- **Neo:** Uses deep RL for complex queries.
- **DQ:** Handles cardinality estimation well.
- **LRO:** Reduces tail latency using real-time feedback.

# Limitations of Related Work



Most learned optimizers suffer from limitations which were not discussed until the next paper came out and addressed them.

# Limitations of Related Work



Most learned optimizers suffer from limitations which were not discussed until the next paper came out and addressed them.

- **Bao:** Enhances optimizers with ML-based query hints, but limits the search space.
- **Neo:** Uses deep RL for complex queries, but suffers from long convergence times.
- **DQ:** Handles cardinality estimation well, but cannot handle sub-nested queries.
- **LRO:** Reduces tail latency using real-time feedback, but requires tight integration with query execution loops.

# Limitations of Related Work



Most learned optimizers suffer from limitations which were not discussed until the next paper came out and addressed them.

- **Bao:** Enhances optimizers with ML-based query hints, but limits the search space.
- **Neo:** Uses deep RL for complex queries, but suffers from long convergence times.
- **DQ:** Handles cardinality estimation well, but cannot handle sub-nested queries.
- **LRO:** Reduces tail latency using real-time feedback, but requires tight integration with query execution loops.


## Cardinality Estimation:

**DeepDB** Integrates learned models for better accuracy in skewed data scenarios.

## Workload Adaptivity:

**Leo** uses feedback for workload shifts.

# How do we evaluate the robustness of a query optimizer?



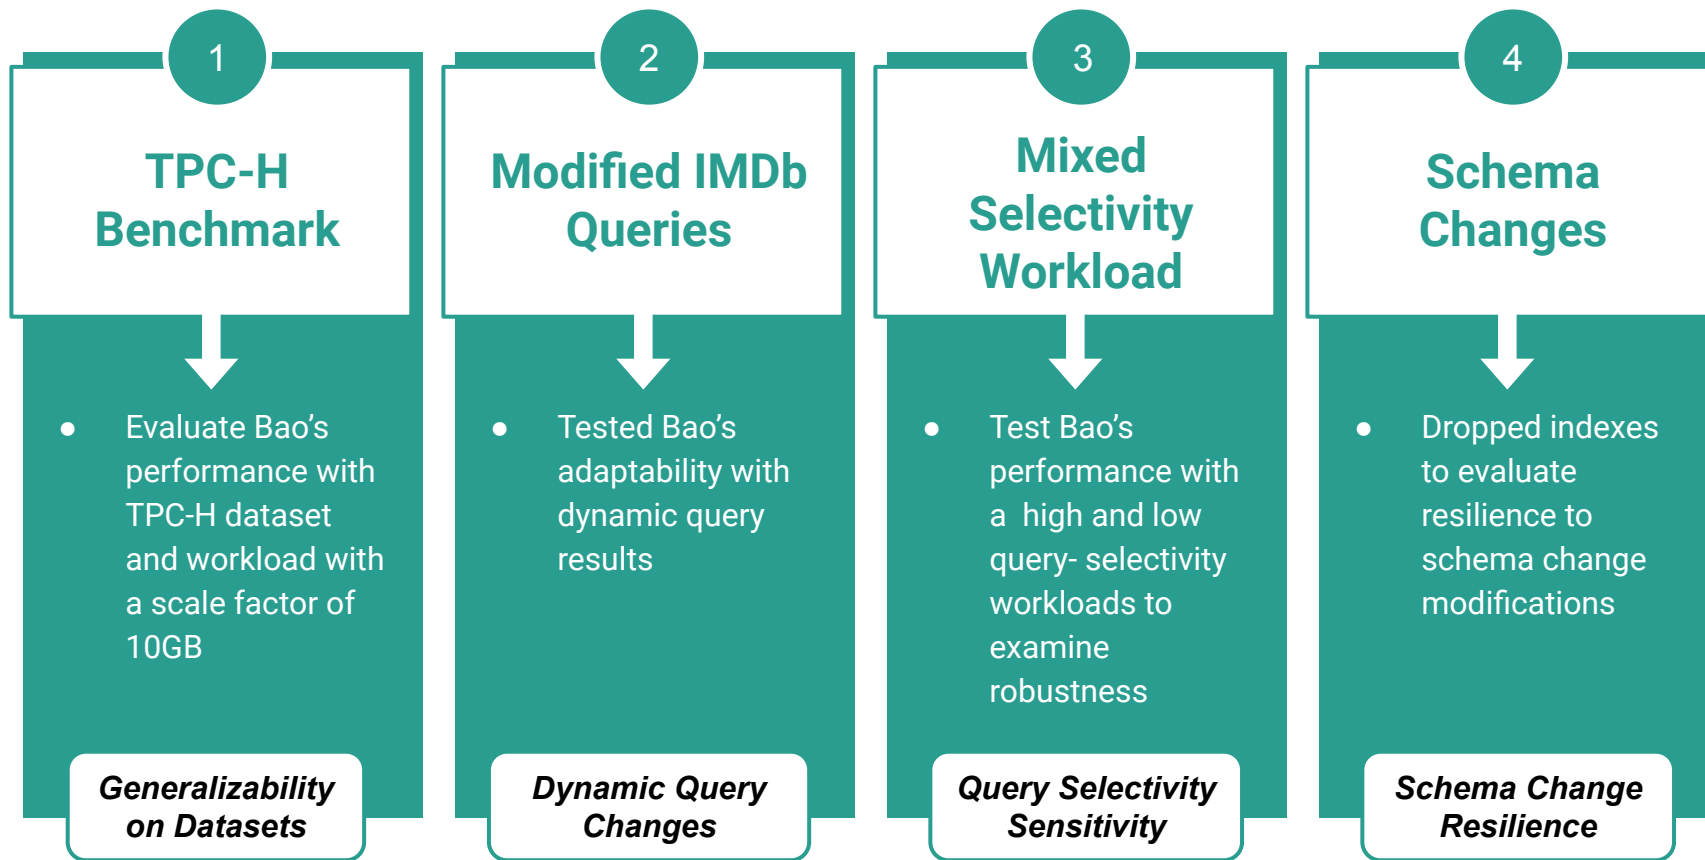
## **Breaking Bao (et. al.) with Systematic Stress Testing**

Designed novel experiments to identify what are the boundaries of Bao's performance, providing insights into its assumptions and limitations.





# Novel Breaking Methodology





# Empirical Evaluation: TPC-H Benchmark

1

## TPC-H Benchmark

- Evaluate Bao's performance with TPC-H dataset and workload with a scale factor of 10GB

**Generalizability  
on Datasets**



Bao's performance was at-par with what was observed in the paper

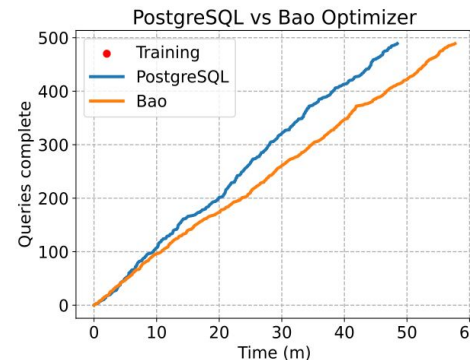



# Empirical Evaluation: Modified IMDb queries

2

## Modified IMDb Queries

- Tested Bao's adaptability with dynamic query results

**Dynamic Query Changes**



**After Modification**

Bao struggled when it encountered the same query with a slight modification in query filter attributes

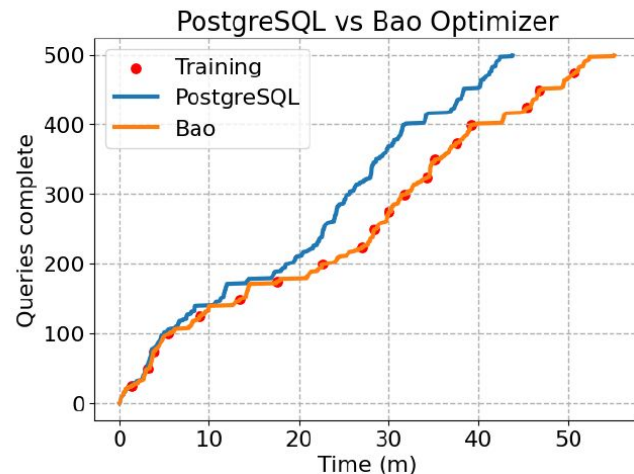

# Empirical Evaluation: Mixed Selectivity Workload

3

## Mixed Selectivity Workload

- Test Bao's performance with a mixed query-selectivity workload to examine robustness

*Mixed Query Workload*



Bao's performance was similar to that of Postgres suggesting that Bao's learning did not add any value.



# Empirical Evaluation: Schema Change

4

## Schema Changes

- Dropped indexes to evaluate resilience to schema change modifications

**Schema Change Resilience**



Bao's performance remained on par, demonstrating resilience to schema changes.

# Empirical Evaluation: Summary



1

**TPC-H  
Benchmark**



***Generalizability  
on Datasets***

2

**Modified IMDb  
Queries**



***Dynamic Query  
Changes***

3

**Mixed  
Selectivity  
Workload**



***Query Selectivity  
Sensitivity***

4

**Schema  
Changes**



***Schema Change  
Resilience***

## Conclusion - Key Takeaways



- Query optimizers like Bao, Neo, DQ are promising and essential.
- However, their benchmarks should include a more robust set of scenarios to highlight their limitations and assumptions.  
For example: schema changes, query diversity.
- This robust benchmarking would give a better understanding and improve adoption of these systems.

# Thank you!



Reference: Gilligan, V. et. al. **Breaking Bad** [TV series]. Sony Pictures Television.